

Package Managers, through a lens or two

Jake Trevor

If you want to follow along

Slides are here:

Package managers are Ubiquitous

How many can you think of?

Package managers are Ubiquitous

Some are familiar:

- utilities from the system PM (apt, brew, yum ...)
- language specific dependencies (cargo, npm, pip, etc.)

Some are a little more exotic:

- ready-made actions in your CI (github actions)
- docker images on GHCR
- games on your iPad's app store
- minecraft mods via a mod manager (other games available)
- music from spotify

Hygiene factors

- Despite how ubiquitous PMs are, no one spends any time thinking about them
- ... because if you are thinking about your package manager, something has gone wrong.
- A good package manager is a *hygiene factor* - you notice it only by its absence

No one notices a clean bathroom

Similarly, no one ever wants to think about pip.

What is a package manager?

There are three characteristic problems package managers solve:

1. **Install** packages
2. Handle **dependencies**
3. (or perhaps, 2.a) Manage **versions**

Almost no one gets #1 wrong

But the rest is a lottery!

What is a good package manager?

How should a well behaved package manager behave?

Idea: Lenses

Lenses

A lens $\mathcal{L} \Gamma \Delta$ on two types Γ, Δ is a pair of functions

$$\text{Get } (\downarrow) : \Gamma \rightarrow \Delta \quad \text{Put } (\uparrow) : \Delta \times \Gamma \rightarrow \Gamma$$

Γ is the source type - it represents '*the truth*'

Δ is the target type - it represents a *view* of the truth

Lens Laws:

A lens is *well behaved* if it obeys the lens laws:

$$\forall (i : \Gamma) (x : \Delta),$$

put-get	get-put
$\downarrow (\uparrow x i) = x$	$\uparrow (\downarrow i) i = i$

and *very well behaved* if it is idempotent:

put-put	
$\uparrow x (\uparrow x i) = \uparrow x i$	

Lenses?

Morally, this seems like it might be a good model ...

- We have some notion of *the truth* (Γ) (the remote package repository)
- We want a *view of the truth* (Δ) (our local installation of packages)

but the devil is in the details

Package lenses:

Let's say that we have:

- a type P of package names
- and for each package $p:P$, a body of type B

A package manager is then a collection of lenses.

Package lenses:

The source type is a map of names to bodies:

$$\Gamma = P \rightarrow B$$

For each subset of package names $P' \subseteq P$, we have a target type, which is a map from names to bodies:

$$\Delta_{P'} = P' \rightarrow B$$

Package lenses:

In lean, we set this up using the *characteristic predicate* of the subset:

```
-- given h : P -> Prop  
type P' = { p : P // h p }  
type P' = Subtype h
```

The *package manager* itself can be understood as a function from the choice of packages to the corresponding lens:

$$\text{PLens} : (h : P \rightarrow \text{Prop}) \rightarrow \mathcal{L} \Gamma \Delta_{P'}$$

Package lenses:

In this formalism, we have a pleasing set of subtypes: $P' <: P$ and $\Gamma <: \Delta_{P'}$

So `Get` is just a cast from Γ to $\Delta_{P'}$.

```
get : (P -> Prop) -> (Subtype h -> Prop)
get := fun γ p => γ p
```

Package lenses:

Put *overwrites* Γ_P with values from $\Delta_{P'}$ whenever they exist:

```
put : ((Subtype h -> Prop) × (P -> Prop)) -> (P -> Prop)
put := fun (δ, γ) p =>
  if h p
  then δ p
  else γ p
```

Package lenses

- Models the intuitive behaviour of installing packages
- All the lenses are *very well behaved*
 - I have lean proofs to back this up

Dependencies

The Problem: Dependency war

- Let's say that I use package B
- and B depends on A version 1.
- Let's say I also want to use C
- which depends on A version 2.
- Some combinations of packages are *impossible*
- How should a well behaved package manager deal with this?

Dependency war

Dependencies: Internal

- Often, a package *completely encapsulates* its dependencies.
 - In this case, the dependencies are **internal** to the package.
- We can use the package without knowing about its dependencies

Dependencies: External

Packages can be *leaky*

- What if types from a dependency appear in the public interface of the package?
- Then the user of the package is *exposed* to the dependencies.
- These are called **external** dependencies

To be clear:

- External dependencies make solving dependency war very hard.

My view:

- Dependency war between external dependencies in packages is *not* a technical problem.
- It's fundamental to the implementation of the packages.

Some sets of packages are *incompatible*.

Nothing a package manager can do will solve this.

Aside: Can we fix it?

If you allow importing multiple versions of the same package,
e.g. by naming them differently ...

then you can expose both interfaces without a namespace
collision

but doing so is contrary to the *spirit* of versioning

→ Versions provide different ways of doing the same thing
∴ it's weird to use two versions of the same thing

However, dependency war between internal dependencies should not be an issue.

Solution: Dependency Locality

Local Dependencies:

The key idea is that we can load dependencies in a *localised* way:

1. When we load a dependency, it must not have side effects on some global environment
2. To support this, packages need to support *dependency injection*

Python does not have this facility; packages are referred to by name, as objects existing on the path

This is not a feature of the package manager, but of the underlying language/system

It's not pip's fault

Pip sucks because python sucks

Module systems:

There's an extensive body of work on module systems, such as:

- **MixML** for ML (Rossberg & Dreyer)
- **backpack** for haskell (Kilpatrick & al.)
- Π - a language agnostic module calculus; unpublished
(Florisson & Mycroft)

These are much more complicated than I have time to go into here, and also a little out of scope.

Versions

Solution 1: new version, new name

- e.g. `pkg-A-v1` and `pkg-A-v2` are different packages
- This works very nicely with the theory we presented!
- It's also completely useless
 - in that it fails to model the actual problem
- *Dependency war* is just *namespace collision*
- If we really want to think about versions, we need to treat them as more than just 'new names'

How should versions work?

- Versioning software is a solved problem
→ Just do what git does!
- What are the semantics of git?

VCS semantics

We have:

- S - a type of the *items* being versioned
- $H<S>$ - the type of histories
- Hash - an associated type of *hashes*

VCS semantics

We don't care about the content of the types so long as we can define the following two functions:

$$\text{Commit} : H<S> \times \text{Hash} \times S \rightarrow H<S> \times \text{Hash}$$

$$\text{Checkout} : H<S> \times \text{Hash} \rightarrow S$$

This might look familiar ...

$$\Gamma = H<S> \times \text{Hash} \qquad \qquad \Delta = S$$

$$\text{Get} = \text{Checkout} \qquad \text{Put} = \text{Commit}$$

There's an infinite family of version control lenses (VCLs), constructed by the function:

$$\text{VCL} : (S : \text{Type}) \rightarrow \mathcal{L}(H<S> \times \text{Hash}) S$$

VCS: Is it a lens?

- If I checkout what I just committed, I should get the same thing back:

$$\text{Checkout}(\text{Commit } x i) = x$$

- If I check something out, and try to commit it, nothing will happen

→ since Git will not (by default) allow empty commits.

$$\text{Commit}(\text{Checkout } i) i = i$$

- By the same fact, committing the same thing twice does nothing

$$\text{Commit } x (\text{Commit } x i) = \text{Commit } x i$$

So, in fact, this kind of abstract VCS is a lens, and it's *very well behaved!*

- Formalising this turned out to be hard
- I have a (very ugly) lean definition for VCLs
- But the proofs are not finished

What do we want, morally?

- Let's say that we have a *history of bodies* for each package
- For a package $p : P$, let V_p be the set of its *versions*
→ i.e. hashes in its history
- Then for a set of packages P' , a *choice of versions* $V_{P'}$ is an object with type

$$V_{P'} = (p : P') \rightarrow V_p$$

What do we want, morally?

Our source type should be a *map of histories* of package bodies:

$$\Gamma = P \rightarrow H\langle B \rangle$$

As before, for a choice of packages P' versions $V_{P'}$, we have a target type - a map of names to bodies:

$$\Delta_{V_{P'}} = P' \rightarrow B$$

We want a family of lenses, constructed by a function:

$$\text{mkPkgMan} = P' \times V_{P'} \rightarrow \mathcal{L}(P \rightarrow H\langle B \rangle)(P' \rightarrow B)$$

What do we have?

We can construct a package lens from names to histories

$$\text{PLens ref} : \mathcal{L}(P \rightarrow H\langle B \rangle) (\text{Subtype ref} \rightarrow H\langle B \rangle)$$

We can construct a VCL on those histories:

$$VCL : \mathcal{L}(H\langle B \rangle \times \text{Hash}) B$$

Here are two other lenses:

$$ID : \mathcal{L} A A$$

$$\text{Concentrate} : \mathcal{L}((B \rightarrow C) \times (B \rightarrow D)) (B \rightarrow (C \times D))$$

These are both vwb. (proofs in lean)

Constructions on lenses:

Lenses can be tensored (parallel composition):

$$_ \otimes _ : \mathcal{L} A B \rightarrow \mathcal{L} C D \rightarrow \mathcal{L} (A \times C) (B \times D)$$

Composed (sequential composition):

$$_ \circ _ : \mathcal{L} A B \rightarrow \mathcal{L} B C \rightarrow \mathcal{L} A C$$

We can split a lens over a function in the second argument:

$$_ \text{split} _ : \mathcal{L} A (B \rightarrow C) \rightarrow \mathcal{L} C D \rightarrow \mathcal{L} A (B \rightarrow D)$$

These constructions are *behaviour preserving* - their output is vwb. when the input lenses are.

What we get: Versioned Package Lenses

$$\begin{aligned} \text{VPL ref} = & ((\text{PLens ref} \otimes \text{ID}) \\ & \quad \circ \text{Concentrate}) \\ & \quad \text{split (VCL)} \\ : & (P \rightarrow \text{Prop}) \rightarrow \mathcal{L} \\ & ((P \rightarrow H\langle B \rangle) \times V_{P'}) \\ & (P' \rightarrow B) \end{aligned}$$

Since this is formed only by vwb. lenses and behaviour preserving constructions, the result is also vwb.

What we get isn't quite what we wanted

$$\mathcal{L} ((P \rightarrow H\langle B \rangle) \times V_{P'}) (P' \rightarrow B)$$

The target type is correct

$$\Delta = (P' \rightarrow B)$$

but the source type is a little different:

Wanted:

$$\Gamma = P \rightarrow H\langle B \rangle$$

Got:

$$\Gamma = (P \rightarrow H\langle B \rangle) \times V_{P'}$$

The difference

The source type includes our choice of versions $V_{P'}$

Intuitively: Put can change our choice of versions

- From a certain point of view, this makes sense
 - The version you just pushed is not the one you checked out
- From another, it does not:
 - If versions can change with a push, why not the set of packages?
 - i.e. what's the story for 'pushing a new package'?

It may well be that a different approach might yield a more pleasing result

wherein versions are names are treated more similarly.

Food for thought

The wrinkles

- Since the formalisation of VCL is not done, I don't have a verified proof that this is all OK just yet.
- At present, while 'tensoring' is morally right, it's not sufficient.
- We need a dependent tensor - where the second argument can depend on the first

Take care!

Notice that for us, the package registry is a *map of histories*:

$$P \rightarrow H$$

But Nix instead has a *history of maps*:

$$H<P \rightarrow B>$$

nixpkgs (the central repository) is a versioned map of names to recipes

Which is the right way around?

My view: Nix does it wrong

With Nix' system, we refer to packages in a particular version of nixpkgs

Problem 1: bumping a version might not change your installation at all

→ i.e. if the packages were not changed in that commit

Problem 2: Versioned versions

Versioned Versions

Nix allows you to refer to different nixpkgs versions
to allow you to refer to specific builds of a given package

But this can be cumbersome

∴ different versions of a package are often distributed as
different packages

e.g. php82 and php83 are different packages.

php82 might exist in 2 commits

You can refer to *different versions of php82*

→ Ridiculous

Conclusions

- We can model package managers with lenses
 - Lots of the nuance is nicely captured
- Particularly, simple unversioned package managers
- For good dependency behaviour, the underlying system needs a good module system
 - Even if we get everything else right, we can still get trapped by non-local dependencies
- This formalism has some weaknesses, space for development
 - A more equal treatment of versions and package names would be nice

Future work?

- What does proper behaviour for **shared dependencies** look like?
 - What's the next best thing after full locality?
 - What about version constraints?
- **Extra features**
 - e.g. declarative/imperative interfaces, reversibility and uninstallation
 - some of this is already captured by the models I've already presented
- **Security**
 - Push should fail sometimes - is that worth modelling?

Some references:

- [1] Florisson & Mycroft, *Towards a Theory of Packages*
unpublished draft (2016)
- [2] Rossberg & Dreyer, *Mixin' Up the ML Module System* (2013)
- [3] Kilpatrick & al, *Backpack: Retrofitting Haskell with Interfaces* (2014)